

# Homework 6

Student: Nguyen Van Thanh  
ID: 20242084

I use python to solve this problem for my convinience.

Declaration: For some translation from matlab code to python, I do use AI tool for supports.

## 1.

- Downscale camera matrix  $K$ :

The downscaled camera matrix  $K$  is calculated with the fomular in page 13 of the slide:

$$K = \begin{pmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix} \longrightarrow K_{\frac{1}{2}} = \begin{pmatrix} \frac{f_x}{2} & 0 & \frac{c_x+0.5}{2} - 0.5 \\ 0 & \frac{f_y}{2} & \frac{c_y+0.5}{2} - 0.5 \\ 0 & 0 & 1 \end{pmatrix}$$

Therefore, we implement the following code in the downscale function:

```
# --- Downscale camera intrinsics ---
# fomular from slide 13
Kd = K.copy()
Kd[0, 0] *= 0.5 # fx/2
Kd[1, 1] *= 0.5 # fy/2
Kd[0, 2] = 0.5 * (Kd[0, 2]+0.5) - 0.5 # 0.5*(cx+0.5)-0.5
Kd[1, 2] = 0.5 * (Kd[1, 2]+0.5) - 0.5 # 0.5*(cy+0.5)-0.5
```

- Downscale intensity image

We use equation (1) to downscale the intensity image. Equation (1) vectorized as below:

```
# --- Downscale intensity image
H,W = I.shape[0], I.shape[1] # resolution of the input image
# New resolution
H2 = H // 2
W2 = W // 2
# Initialize downsampled intensity image
Id = np.zeros((H2, W2), dtype=I.dtype)
# Vectorized from equation (1)
Id = 0.25 * (
    I[0:H:2, 0:W:2] +
    I[1:H:2, 0:W:2] +
    I[0:H:2, 1:W:2] +
    I[1:H:2, 1:W:2]
)
```

- Downscale Depth Image

We use equation (2) to downscale the depth image. The implementation is as follows:

```
# --- Downscale depth image
# initialize downsampled depth image
Dd = np.zeros((H2, W2), dtype=D.dtype)

# Extract 2x2 blocks, this is D(x,y) set
d00 = D[0:H:2, 0:W:2]
d10 = D[1:H:2, 0:W:2]
d01 = D[0:H:2, 1:W:2]
d11 = D[1:H:2, 1:W:2]

# Stack into (H2, W2, 4), then each pixel has 4 depth values
blocks = np.stack([d00, d10, d01, d11], axis=-1)

# Valid depth mask (non-zero) (H2, W2, 4) shape
valid_mask = blocks != 0 # check which pixel has valid depth, returns boolean (0/1) array

# Sum of valid depth values
valid_sum = np.sum(blocks * valid_mask, axis=-1)
```

```

# Count of valid values (H2, W2) shape
valid_count = np.sum(valid_mask, axis=-1)

# Avoid division by zero → depth stays 0
nonzero = valid_count > 0

# Take average of valid depth values as in equation (2)
Dd[nonzero] = valid_sum[nonzero] / valid_count[nonzero]

# --- Recursive call ---
return downscale(Id, Dd, Kd, level - 1)

```

We test the implemented function with a provided RGB and depth images. Figures below illustrate the downscaled images with different level.



Figure 1: Downscaled RGB images

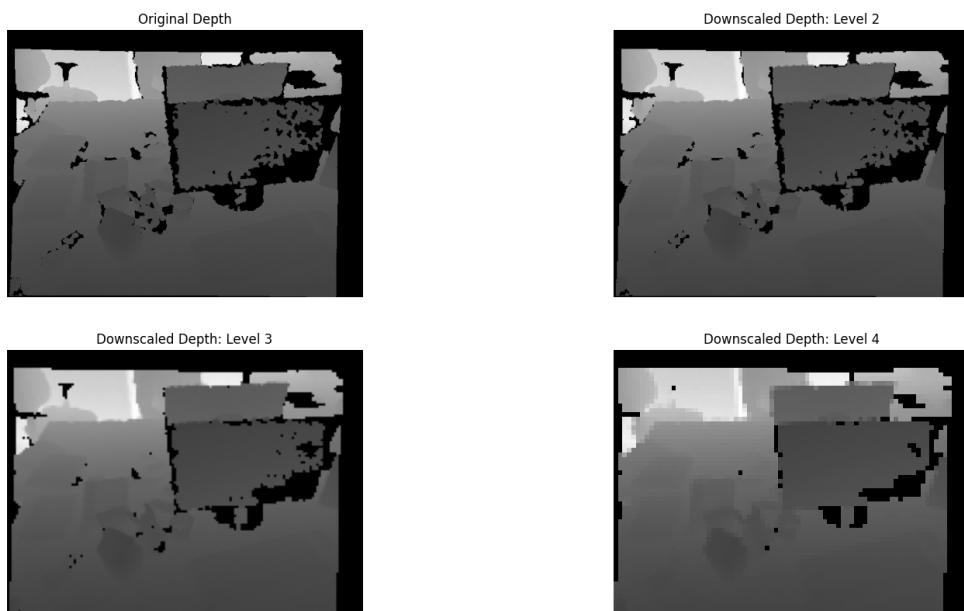


Figure 2: Downscale depth images

## 2.

In order to calculate the per-pixel residual  $\mathbf{r}(\xi)$ , we have two main computations. First, we warps a pixel from a ref.image to new image with  $\omega(\mathbf{p}_i, d, \xi)$ :

$$\omega(\mathbf{p}_i, d, \xi) = \pi(K(R_\xi)K^{-1} \begin{pmatrix} dp_{i,x} \\ dp_{i,y} \\ d \end{pmatrix} + \mathbf{t}_\xi))$$

Then projection on a new image is given by:

$$\pi(x, y, z) = \begin{pmatrix} x/z \\ y/z \end{pmatrix}$$

Then we have the following snippet code

```
# TODO warp reference points to target image
d = DRef[y, x]
if d == 0:
    continue; # skip invalid depth
# backproject to 3D point (3,1)
temp = K @ (R @ K_inv @ np.array([x, y, 1]) * d + t)
w = np.array([temp[0] / temp[2], temp[1] / temp[2]]) # [x/z, y/z]
# TODO project warped points onto image
xImg[y, x] = w[0]
yImg[y, x] = w[1]
```

Then the 2nd image is used to interpolate the intensity value given the calculated coordinates to archive new image and calculate the residual finally:

```
coords = np.array([yImg.ravel(), xImg.ravel()])
I_warped = map_coordinates(I, coords, order=1, mode='nearest')
I_warped = I_warped.reshape(H_ref, W_ref)
# --- 5. Photometric residual
err = IRef - I_warped
return err.reshape(-1, 1)
```

The residual contains intensity values and we can visualize as an image as below:

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from downscale import downscale
# use cv2.imread to read grayscale images as float
# because map_coordinates requires float inputs
IRef = cv2.imread("rgb/1305031102.175304.png", cv2.IMREAD_GRAYSCALE).astype(float)
DRef = cv2.imread("depth/1305031102.160407.png", cv2.IMREAD_GRAYSCALE).astype(float)
I = cv2.imread("rgb/1305031102.275326.png", cv2.IMREAD_GRAYSCALE).astype(float) # new image
# camera intrinsics
K = np.array([[525.0, 0.0, 319.5],
              [0.0, 525.0, 239.5],
              [0.0, 0.0, 1.0]])
# camera pose 6D vector
xi = np.array([0.0, 0.0, 0.0, 0.0, 0.0, 0.0]) # small rotation around z-axis
err = calcErr(IRef, DRef, I, xi, K)
# visualize error as image
err_image = err.reshape(IRef.shape)
plt.imshow(err_image, cmap='gray')
plt.title('Photometric Error Image')
plt.axis('off')
plt.show()
```

We have the following figure

Photometric Error Image



Figure 3: A calculated residual image

### 3.

The function to calculate the numeric derivative of the residual is implemented as belows:

```

1 def deriveErrNumeric(IRef, DRef, I, xi, K):
2     # calculate numeric derivative (SLOW!!!)
3
4     # compute residuals for xi
5     residual_xi = calcErr(IRef,DRef,I,xi,K)
6
7     H_I, W_I = I.shape[0], IRef.shape[1] # resolution of the reference image, heght and width
8     num_of_pixels = H_I * W_I
9     # initialize Jacobian
10    Jac = np.zeros((num_of_pixels, 6))
11
12    # compute Jacobian numerically
13    eps = 1e-6
14    for j in range(1, 7):
15        epsVec = np.zeros((6, 1))
16        epsVec[j-1] = eps
17        # multiply epsilon from left onto the current estimate.
18        xiPerm = se3Log(se3Exp(epsVec) @ se3Exp(xi))
19        # TODO compute respective column of the Jacobian
20        # (hint: difference between residuals of xi and xiPerm)
21        residual_perm = calcErr(IRef, DRef, I, xiPerm, K)
22        Jac[:, j-1] = (residual_perm - residual_xi).flatten() / eps
23    return Jac, residual_xi

```

We use the calcErr() function implemented previously. Here we explain the main calculation in the function. The following line

```
1 xiPerm = se3Log(se3Exp(epsVec) @ se3Exp(xi))
```

corresponds to:

$$\xi_{\text{permutation}} = \epsilon \mathbf{e}_1 \mathbf{o} \xi$$

and

```
1 residual_perm = calcErr(IRef, DRef, I, xiPerm, K)
```

calculate  $r(\epsilon \mathbf{e}_1 \mathbf{o} \xi)$ . Finally, the corresponding column of the Jacobian matrix is calculated as:

```
1 Jac[:, j-1] = (residual_perm - residual_xi).flatten() / eps
```

which corresponds to:

$$J = \frac{r(\epsilon e_1 o \xi) - r(\xi)}{\epsilon}$$

## 4.

The implementation of the Gauss-Newton step is referred to slide number 10, particularly on the calculation of the updates  $\delta_\xi$ :

$$\delta_\xi = -(J_r^T J_r)^{-1} J_r^T r_0$$

and then update the solution:

$$\xi^{k+1} = \delta_\xi o \xi^k$$

Those calculations are captured in the snip code below:

```
1 # -----
2 # Gauss-Newton step
3 # -----
4 H = Jac.T @ Jac
5 b = Jac.T @ residual
6 delta_xi = -np.linalg.solve(H, b) # solve for delta
7 # -----
8 # delta_xiate pose
9 # xi = se3Log(se3Exp(delta_xi) * se3Exp(xi))
10 # -----
11 lastXi = xi.copy()
12 # update solution
13 xi = se3Log(se3Exp(delta_xi.flatten()) @ se3Exp(xi.flatten())).reshape(6, 1)
```

The capture below shows the printed result when we run the doAlignment program:

```
(computer vision) vanthanhanguyen@vanthanhanguyen:~/Documents/Documents/Multiple_Geometry_for_Spatial_AI/HW6$ /usr/bin/env
/home/vanthanhnguyen/miniconda3/envs/computer_vision/bin/python /home/vanthanhnguyen/.vscode/extensions/ms-python.debugpy-
2025.16.0-linux-x64/bundled/libs/debugpy/adapter/.../debugpy/launcher 38205 -- /home/vanthanhnguyen/Documents/Documents/
Multiple_Geometry_for_Spatial_AI/HW6/doAlignment.py
xi = [-0.00019604  0.00476781  0.03822938 -0.02964087 -0.01937721 -0.00063463]
error = 1847.0283264624754
Level: 4
xi = [-0.00185411  0.00525217  0.03692122 -0.02939112 -0.01851548 -0.00179947]
error = 1570.50602632311
xi = [-0.0017135   0.00514199  0.03684258 -0.0294524  -0.01860742 -0.00182484]
error = 1568.492553766944
Level: 3
xi = [-0.00233905  0.00498869  0.03800416 -0.02966783 -0.01827206 -0.00168796]
error = 1443.49578535251
xi = [-0.00225367  0.0049256   0.03802295 -0.02969241 -0.01833218 -0.00170998]
error = 1442.3234632875333
Level: 2
xi = [-0.0024604   0.0052894   0.03761545 -0.02946776 -0.01812504 -0.00144994]
error = 1812.4787432413393
xi = [-0.00246431  0.00542533  0.03753814 -0.0293851  -0.0180866  -0.00137731]
error = 1811.6393980481787
Level: 1
xi = [-0.00229223  0.00581735  0.03748227 -0.02911642 -0.01816053 -0.00104458]
error = 1954.3007317854212
xi = [-0.00221624  0.00595091  0.03748999 -0.02902572 -0.01821022 -0.00096311]
error = 1953.2514422208164
```

Figure 4: Estimation result of the campera pose

We see that the final result for the level-1 image is:

$$\xi = (-0.00221624 \quad 0.00595091 \quad 0.03748999 \quad -0.02902572 \quad -0.01821022 \quad -0.00096311)$$

## 5.

We briefly explain the calculation procedure for the analytical Jacobian.

- Calculate image gradient

Given a pixel at  $(x, y)$ , the image gradient at that point is calculated as follows:

$$\begin{aligned}\frac{\partial I}{\partial x}(x, y) &= \frac{I(x+1, y) - I(x-1, y)}{2} \\ \frac{\partial I}{\partial y}(x, y) &= \frac{I(x, y+1) - I(x, y-1)}{2}\end{aligned}$$

Therefore, we have following snip code:

```
(computer_vision) vanthanhanguyen@vanthanhanguyen:~/Documents/Documents/Multiple_Ge  
/home/vanthanhnguyen/miniconda3/envs/computer_vision/bin/python /home/vanthanhng  
2025.16.0-linux-x64/bundled/libs/debugpy/adapter/.../debugpy/launcher 46655 --  
Multiple Geometry for Spacial AI/HW6/doAlignment.py  
Level: 4  
xi = [-0.00373062  0.004184   0.03778246 -0.03036128 -0.01720939 -0.00213818]  
error = 99.18051668725062  
xi = [-0.00277322  0.00496207  0.03867736 -0.02974924 -0.01780257 -0.0017153 ]  
error = 96.876213992142  
xi = [-0.00307249  0.00446159  0.03825833 -0.03013815 -0.01759944 -0.00184856]  
error = 96.49507628168764  
Level: 3  
xi = [-0.00276867  0.00529128  0.03792075 -0.02941678 -0.01799551 -0.001532 ]  
error = 104.9773751105924  
xi = [-0.00300256  0.00512392  0.03813566 -0.02958592 -0.01776379 -0.00173837]  
error = 104.16276311051365  
xi = [-0.00288139  0.00514546  0.03801145 -0.02954262 -0.01787625 -0.00163867]  
error = 104.08594467791848  
Level: 2  
xi = [-0.00249836  0.00537972  0.03755556 -0.02947296 -0.01801245 -0.00116952]  
error = 95.34181081796682  
xi = [-0.00250575  0.00532844  0.03768043 -0.02951618 -0.01803442 -0.00114593]  
error = 94.80692919976153  
xi = [-0.00245812  0.00531973  0.03766641 -0.02952903 -0.01806281 -0.00113571]  
error = 94.89233974874281  
Level: 1  
xi = [-0.00206844  0.00559285  0.03743413 -0.02931957 -0.01829692 -0.00091777]  
error = 110.39652267469175  
xi = [-0.0020301  0.00573002  0.03747901 -0.02921963 -0.01833009 -0.00086705]  
error = 109.32088512442266  
xi = [-0.00201632  0.00576376  0.03748579 -0.02919403 -0.01834147 -0.00085957]  
error = 109.21526260688103
```

Figure 5: Camera estimation result using analytic Jacobian

```

1 # =====
2 # Image gradients (central differences)
3 # =====
4 dxI = np.zeros_like(I)
5 dyI = np.zeros_like(I)
6 dxI[:, 1:-1] = 0.5 * (I[:, 2:] - I[:, :-2]) # exclude 1st and last column
7 dyI[1:-1, :] = 0.5 * (I[2:, :] - I[:-2, :]) # exclude 1st and last row

```

- Calculate Jacobian

The Jacobian is calculated as in slide 11:

$$\left. \frac{\partial r_i(\epsilon \circ \xi^{(k)})}{\partial \epsilon} \right|_{\epsilon=0} = -\frac{1}{z'} (\nabla_{I_x} f_x \quad \nabla_{I_y} f_y) \begin{pmatrix} 1 & 0 & -\frac{x'}{z'} & -\frac{x'y'}{z'} & \left(z' + \frac{x'^2}{z'}\right) & -y' \\ 0 & 1 & -\frac{y'}{z'} & -\left(z' + \frac{y'^2}{z'}\right) & \frac{x'y'}{z'} & x' \end{pmatrix}$$

Therefore, we have following snip code:

```

1 # =====
2 # Analytic Jacobian (Kerl Eq. 4.14)
3 # =====
4 Jac = np.zeros((N, 6))
5
6 inv_z = 1.0 / zp
7 inv_z2 = inv_z ** 2
8
9 Jac[:, 0] = dxInterp * inv_z
10 Jac[:, 1] = dyInterp * inv_z
11 Jac[:, 2] = -(dxInterp * xp + dyInterp * yp) * inv_z2
12
13 Jac[:, 3] = -(dxInterp * xp * yp) * inv_z2 - dyInterp * (1 + (yp * inv_z) ** 2)
14 Jac[:, 4] = dxInterp * (1 + (xp * inv_z) ** 2) + dyInterp * xp * yp * inv_z2
15 Jac[:, 5] = -(dxInterp * yp - dyInterp * xp) * inv_z

```

The final result as we run the doAlignment program is shown as below: As we see, the final estimation is:

$$\xi = \begin{pmatrix} -0.00206844 & 0.00559285 & 0.03743413 & -0.02931957 & -0.01829692 & -0.00091777 \end{pmatrix}$$

## **Appendix**

Source code:

[https://github.com/ThanhNV-Robotics/HW6\\_Multiple\\_Geometry\\_for\\_SpatialAI](https://github.com/ThanhNV-Robotics/HW6_Multiple_Geometry_for_SpatialAI)