

Multiple View Geometry: Exercise #6

Exercise 6.1

In this exercise you will implement direct image alignment as Gauss-Newton minimization on SE(3). Download the package `hw6.zip` provided on the website. It contains a code framework, test images and the corresponding camera calibration.

1. Implement a function `[Id, Dd, Kd] = downscale(I, D, K, level)` which (recursively) halves the image resolution of the image I , the depth map D and adjusts the corresponding camera matrix K per pyramid level (according to the slides). For an input frame of dimensions 640×480 (level 1), level 2 corresponds to 320×240 pixels, level 3 corresponds to 160×120 pixels and so on. For the intensity image, downscaling is performed by averaging the intensity, that is

$$I_d(x, y) := 0.25 \sum_{x', y' \in O(x, y)} I(x', y') \quad (1)$$

where $O(x, y) = \{(2x, 2y), (2x + 1, 2y), (2x, 2y + 1), (2x + 1, 2y + 1)\}$.

For the depth map, downscaling is performed by averaging the depth of all valid pixels (invalid depth values are set to zero), that is

$$D_d(x, y) := \left(\sum_{x', y' \in O_d(x, y)} D(x', y') \right) / |O_d(x, y)| \quad (2)$$

where $O_d(x, y) := \{(x', y') \in O(x, y) : D(x', y') \neq 0\}$.

2. Complete the function `r = calcErr(I1, D1, I2, xi, K)` that takes the images and their (assumed) relative pose, and calculates the per-pixel residual $\mathbf{r}(\xi)$ as defined in the slides. \mathbf{r} should be a $n \times 1$ vector, with $n = w \times h$. Visualize the residual as image for $\xi = 0$.
Hint: work on a coarse version of the image (e.g. 160×120) to make it run faster.
3. Implement a function `[J, r] = deriveNumeric(I1, D1, I2, xi, K)` that **numerically** derives $\mathbf{r}(\xi)$ on the manifold, i.e., for each pixel i computes

$$\frac{\partial r_i(\xi)}{\partial \xi} = \left(\frac{r_i((\epsilon \mathbf{e}_1) \circ \xi) - r_i(\xi)}{\epsilon}, \dots, \frac{r_i((\epsilon \mathbf{e}_6) \circ \xi) - r_i(\xi)}{\epsilon} \right) \quad (3)$$

where ϵ is a small value (for Matlab $\epsilon = 10^{-6}$), and \mathbf{e}_j is the j 'th unit vector. J should be a $n \times 6$ matrix. Additionally, the per-pixel residuals $\mathbf{r}(\xi)$ are returned as r .

4. Implement Gauss Newton minimization for the photometric error $E(\xi) = \|\mathbf{r}(\xi)\|_2^2$ as derived in the slides. Use only one pyramid level (160×120) in the beginning, and then add the others.
5. Implement a function `J = deriveAnalytic(I1, D1, I2, xi, K)` which **analytically** derives $\mathbf{r}(\xi)$ (see slides). Using it instead of the numeric derivatives in the minimization from the previous task should result in a significant speed-up.

Direct Image Alignment

- = “Direct Tracking” / “Dense Tracking” / “Dense Visual Odometry”
- = “Lucas-Kanade Tracking on $\text{SE}(3)$ ”

ref. image



ref. depth



+



new image

->

Camera
pose ξ

Direct Image Alignment

Robust Odometry Estimation for RGB-D Cameras

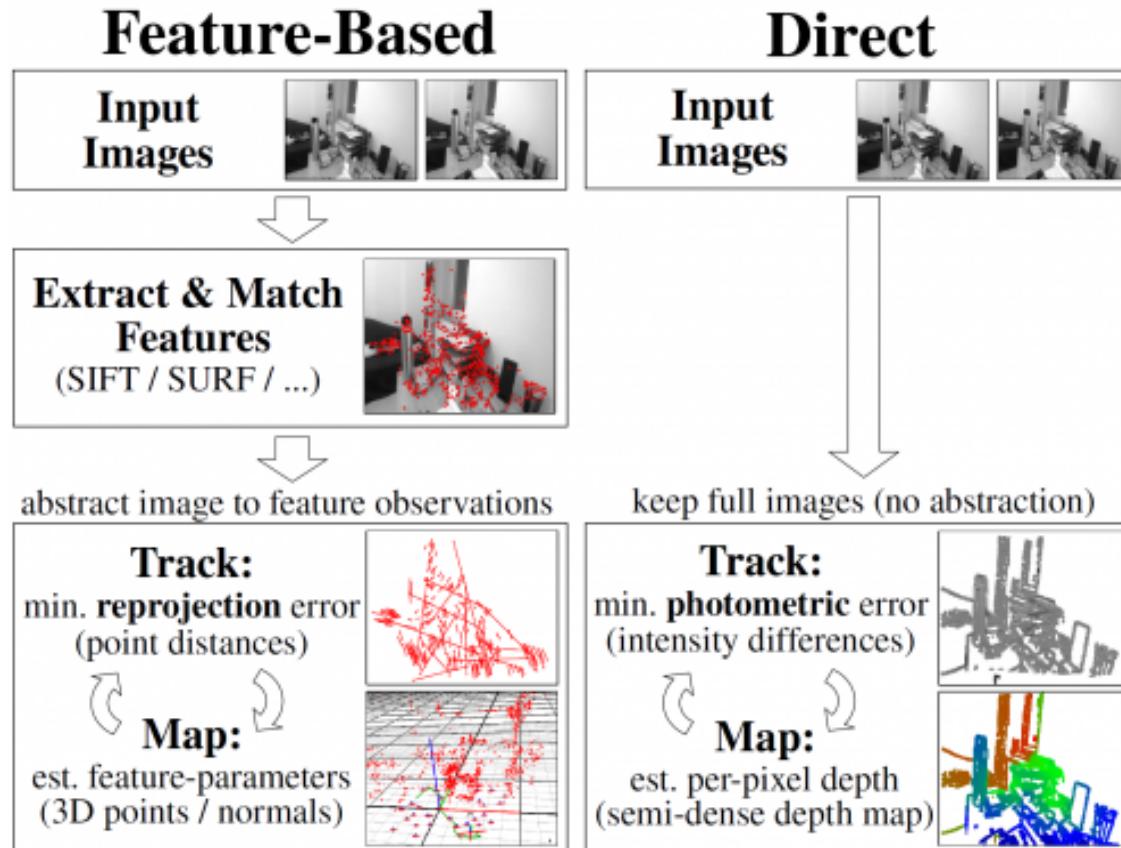
Christian Kerl, Jürgen Sturm, Daniel Cremers



Computer Vision and Pattern Recognition Group
Department of Computer Science
Technical University of Munich

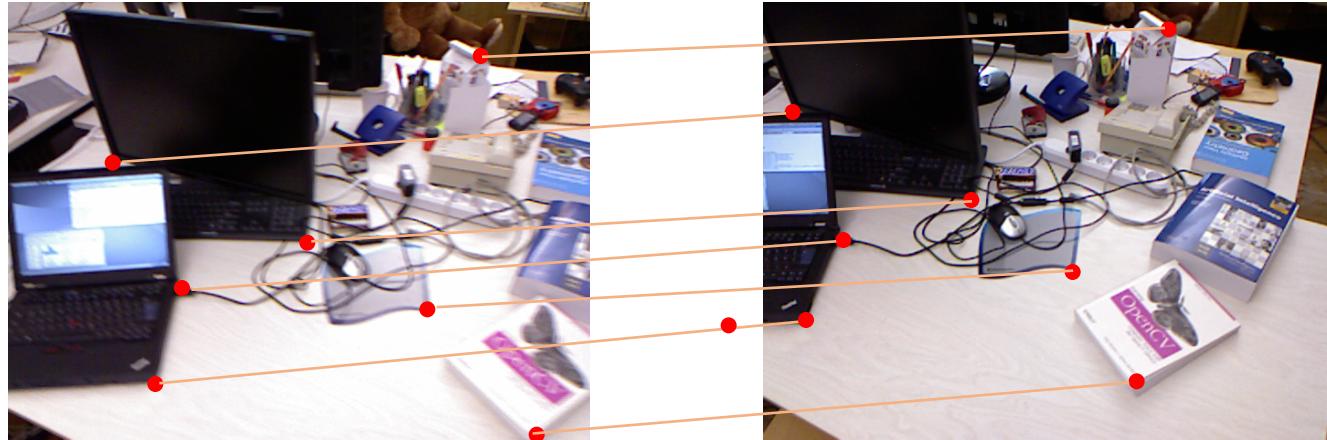


Keypoints, Direct, Sparse, Dense

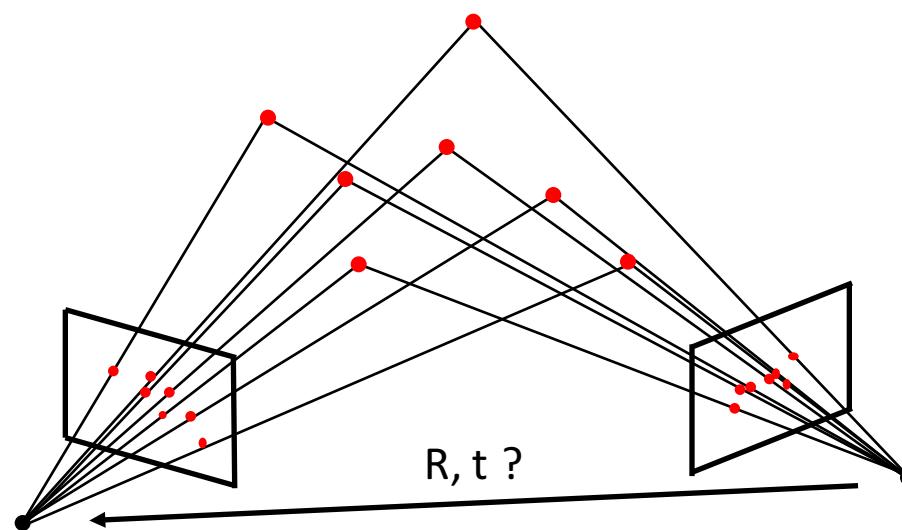


- Sparse: use a small set of selected pixels (keypoints)
- Dense: use all (valid) pixels

Sparse Keypoint-based Visual Odometry

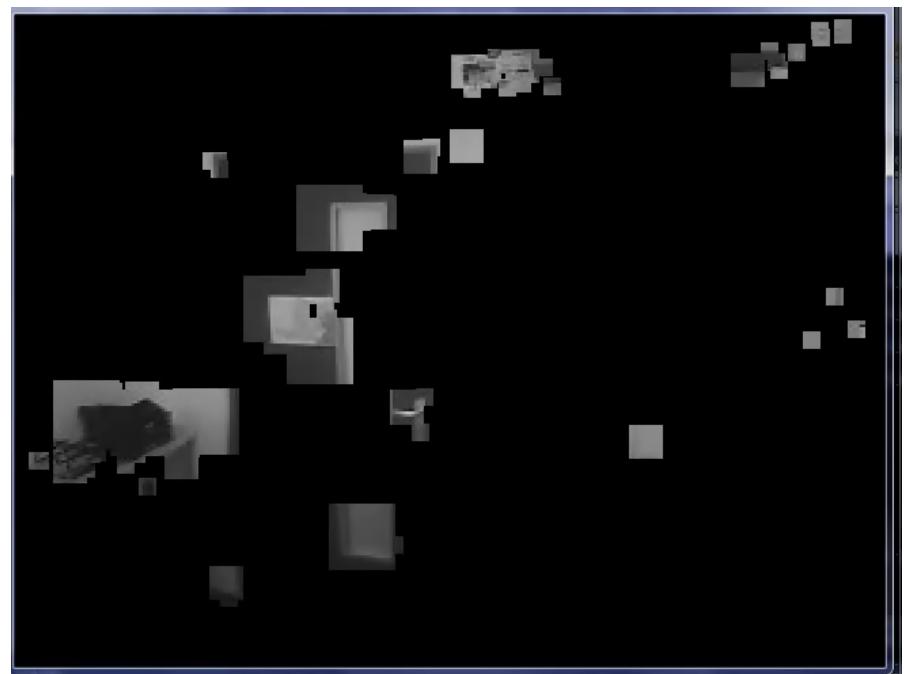


Extract and match
keypoints

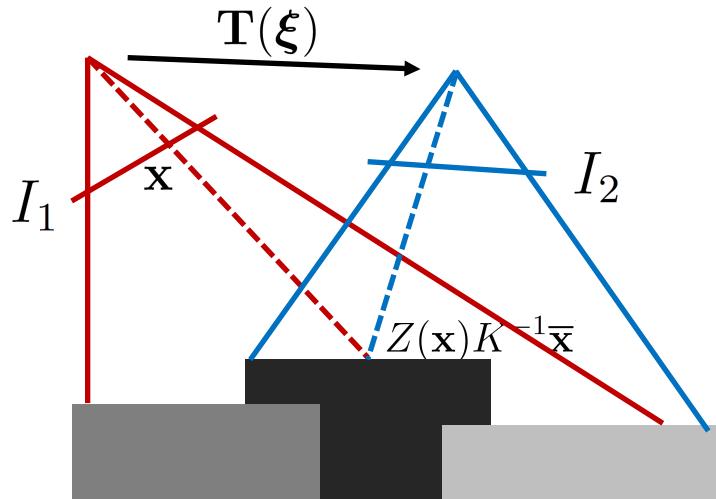


Determine relative
camera pose (R, t)
from keypoint matches

Problem with Keypoint-based Methods



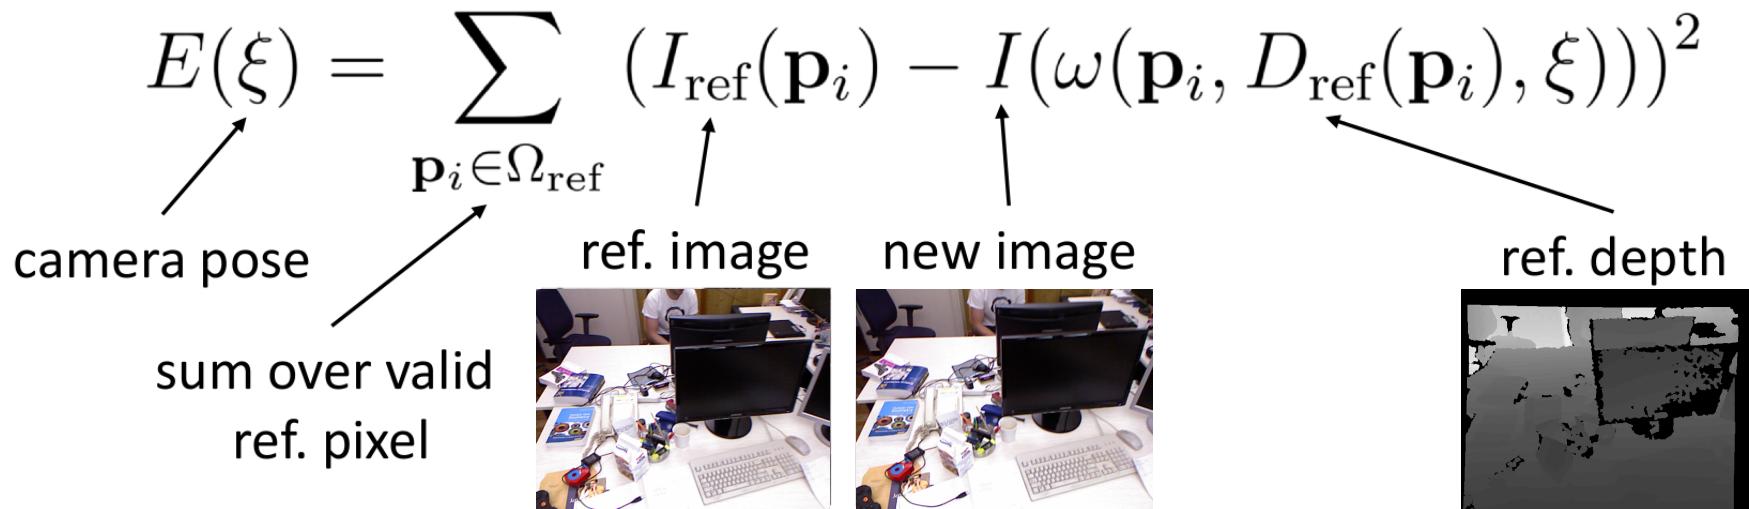
Dense Direct Image Alignment



- Known pixel depth -> "simulate" RGB-D image from different view point
- Ideally: warped image = image taken from that pose:

$$I_1(\mathbf{x}) = I_2(\pi(\mathbf{T}(\xi)Z(\mathbf{x})K^{-1}\bar{\mathbf{x}}))$$
- RGB-D: depth available -> find camera motion!
- Motion representation using the SE(3) Lie algebra
- Non-linear least squares optimization

Direct minimization of photometric error



$$\omega(\mathbf{p}_i, d, \xi) := \pi(K(R_\xi K^{-1} \begin{pmatrix} dp_{i,x} \\ dp_{i,y} \\ d \end{pmatrix} + \mathbf{t}_\xi))$$

$$\pi(x, y, z) := \begin{pmatrix} x/z \\ y/z \end{pmatrix}$$

$$\begin{pmatrix} R_\xi & \mathbf{t}_\xi \\ \mathbf{0} & 1 \end{pmatrix} := \exp(\hat{\xi})$$

$\omega(\mathbf{p}_i, d, \xi)$ "warps" a pixel from
ref. image to new image

Gauss-Newton

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} (I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2$$

- Solved using the **Gauss-Newton** algorithm using left-multiplicative increments on SE(3):

$$\xi_1 \circ \xi_2 := \log(\exp(\hat{\xi}_1) \cdot \exp(\hat{\xi}_2))^\vee \neq \xi_1 + \xi_2$$

$$\neq \xi_2 \circ \xi_1$$

- Intuition: iteratively solve for $\nabla E(\xi) = 0$ by approximating $\nabla E(\xi)$ linearly (i.e. by approximating $E(\xi)$ quadratically)
- Using **coarse-to-fine** pyramid approach

Gauss-Newton

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} \underbrace{(I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2}_{=: r_i^2(\xi)}$$

1. „Linearize“ \mathbf{r} on Manifold around current pose $\xi^{(n)}$:

$$\mathbf{r}(\xi) \approx \underbrace{\mathbf{r}(\xi^{(k)})}_{\mathbf{r}_0 \in R^n} + \underbrace{\frac{\partial \mathbf{r}(\epsilon \circ \xi^{(k)})}{\partial \epsilon}}_{J_{\mathbf{r}} \in R^{n \times 6}} \Big|_{\epsilon=0} \cdot \underbrace{(\xi \circ (\xi^{(k)})^{-1})}_{\delta_{\xi}}$$

2. Solve for $\nabla E(\xi) = 0$

$$E(\xi) = \|\mathbf{r}_0 + J_{\mathbf{r}} \cdot \delta_{\xi}\|_2^2 = \mathbf{r}_0^T \mathbf{r}_0 + 2\delta_{\xi}^T J_{\mathbf{r}}^T \mathbf{r}_0 + \delta_{\xi}^T J_{\mathbf{r}}^T J_{\mathbf{r}} \delta_{\xi}$$

$$\nabla E(\xi) = 2J_{\mathbf{r}}^T \mathbf{r}_0 + 2J_{\mathbf{r}}^T J_{\mathbf{r}} \delta_{\xi} = 0 \quad \Rightarrow \quad \delta_{\xi} = -(J_{\mathbf{r}}^T J_{\mathbf{r}})^{-1} J_{\mathbf{r}}^T \mathbf{r}_0$$

3. Apply $\xi^{(k+1)} = \delta_{\xi} \circ \xi^{(k)}$

4. Iterate (until convergence)



Gauss-Newton

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} \underbrace{(I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2}_{=: r_i^2(\xi)}$$

- Jacobian J_r : partial derivatives

Requires gradient of residual:

$$\frac{\partial r_i(\epsilon \circ \xi^{(k)})}{\partial \epsilon} \Big|_{\epsilon=0} = -\frac{1}{z'} \begin{pmatrix} \nabla I_x f_x & \nabla I_y f_y \end{pmatrix} \begin{pmatrix} 1 & 0 & -\frac{x'}{z'} & -\frac{x'y'}{z'} & \left(z' + \frac{x'^2}{z'}\right) & -y' \\ 0 & 1 & -\frac{y'}{z'} & -(z' + \frac{y'^2}{z'}) & \frac{x'y'}{z'} & x' \end{pmatrix}$$

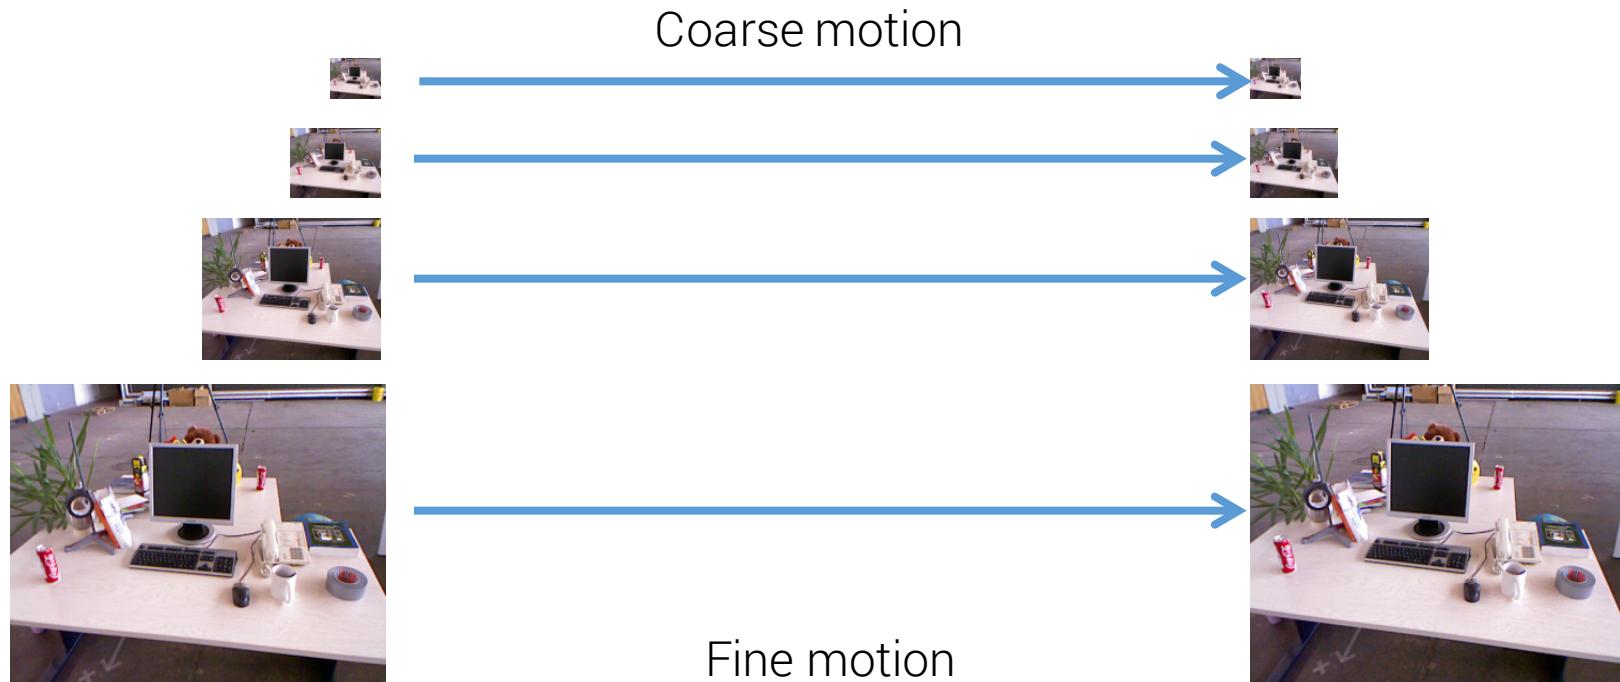
$= 1 \times 6 \text{ row of } J_r$

with

- $\begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} := R_{\xi^{(k)}} K^{-1} \begin{pmatrix} dp_{i,x} \\ dp_{i,y} \\ d \end{pmatrix} + \mathbf{t}_{\xi^{(k)}}$ = warped point (before projection)
- f_x, f_y, K = intrinsic camera calibration
- $\nabla I_x, \nabla I_y$ = image gradients

Coarse-to-Fine

- Adapt size of the neighborhood from coarse to fine



Coarse-to-Fine

- Minimize for down-scaled image (e.g. factor 8, 4, 2, 1) and use result as initialization for next finer level
- Elegant formulation: Downscale image and adjust K accordingly
 - Downscale by factor of 2 (e.g. 640x480 -> 320x240)
 - Camera matrix

$$K = \begin{pmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix} \Rightarrow K_{\frac{1}{2}} = \begin{pmatrix} \frac{f_x}{2} & 0 & \frac{c_x+0.5}{2} - 0.5 \\ 0 & \frac{f_y}{2} & \frac{c_y+0.5}{2} - 0.5 \\ 0 & 0 & 1 \end{pmatrix}$$

(assuming discrete pixel (x,y) contains continuous value at (x,y))



Final Algorithm

$$\xi^{(0)} = \mathbf{0}$$

k = 0

for *level* = 3 ... 1

 compute down-scaled images & depthmaps (factor = 2^{level})

 compute down-scaled K (factor = 2^{level})

for *i* = 1..20

 compute Jacobian $J_r \in R^{n \times 6}$

 compute update δ_ξ

 apply update $\xi^{(k+1)} = \delta_\xi \circ \xi^{(k)}$

k++; maybe break early if δ_ξ too small or if residual increased

done

done

+ robust weights (e.g. Huber), see *iteratively reweighted least squares*

+ *Levenberg-Marquard (LM)* Algorithm